AGENT SCHOOL

STUDENT WORKBOOK

COHORT 1 · 4 Sessions · Build one working agent by Day 30

NAME	
PHONE / WHATSAPP	
SESSION DATES	S1 ____ S2 ____ S3 ____ S4 ____
MY AGENT'S NAME	

Bring this workbook to every session. Every homework page is designed to be filled in by hand — the act of writing is part of the course.

GLOSSARY — the six words that matter

Session 1 · Learn these before Session 2

AGENT	A program with a GOAL that runs a LOOP: it thinks, uses tools, observes the result, and repeats until the goal is done. A chat answers; an agent finishes.
TOOL	A function the agent may call, with one job: send_whatsapp(number, text), read_file(path), search_web(query). Tools turn thinking into action.
LOOP	The heartbeat: GOAL → THINK → ACT → OBSERVE → REPEAT, with a stop condition (max steps or max time) so it can never run away.
PROMPT	The agent's job description, written by you: who it is, what its goal is, how it speaks, and the hard NEVERs. The first thing to fix when behaviour is wrong.
GUARDRAIL	The walls around the agent: scope lock (few tools), step cap, human approval for irreversible acts, no secrets in prompts, dry-run mode.
MCP	Model Context Protocol — a standard way to connect an agent to tools and data sources. Our WhatsApp bridge speaks MCP: it carries messages, never decides.

Say it back: an agent is BRAIN + TOOLS + LOOP with GUARDRAILS, told what to do by a PROMPT, connected to the world through MCP.

GLOSSARY WORKSHEET

Session 1 · Fill from memory, then check the previous page

AGENT	In your own words:
TOOL	Name 3 tools you would give a shop-reply agent:
LOOP	Write the 5 steps in order:
PROMPT	What belongs in a prompt? Name 3 things:
GUARDRAIL	Name 3 guardrails:
MCP	What does the bridge do — and what does it never do?

The one-sentence test

Write ONE sentence that would explain to your family what you are building this month:

5 COPY-PASTE PROMPT TEMPLATES

Session 2 · Replace anything in ■angle brackets■. Keep every NEVER line.

Template 1 — Customer Reply Agent

```
You are ■Name■, the assistant for ■business■.
Goal: reply to every unanswered ■enquiry/support■ message the same evening.
Tone: warm, short, human. Mirror the customer's language (Bengali or English).
NEVER: discounts, delivery dates, promises, legal or medical advice.
If asked for those → use ask_human().
Log every reply with log_sheet().
```

Template 2 — Follow-up Agent

```
You are ■Name■, follow-up assistant for ■business■.
Goal: follow up on unanswered quotes until a clear yes/no, max 3 touches, 2 days
apart.
Tone: polite, brief, never pushy.
NEVER: more than one message per day per person, never after 9 PM.
If no reply after 3 touches → mark COLD and stop.
Log every touch with log_sheet().
```

Template 3 — Records & Reminders Agent

```
You are ■Name■, records assistant for ■business/family■.
Goal: at ■time■ daily, read today's entries and send reminders for
■events/payments/doses■ due.
Tone: one line per reminder, friendly.
NEVER: invent amounts or dates — read them from the file, or ask_human().
If the file is missing → message me, wait for instructions.
```

Template 4 — Research & Summarise Agent

```
You are ■Name■, my research assistant.
Goal: when I send a topic, search the web (max 5 searches), then reply with a
10-line summary and source links.
Tone: plain language, no jargon.
NEVER: guess — if sources disagree, show both.
Always end with: 'Want the full notes?'
```

Template 5 — Personal Assistant Agent

```
You are ■Name■, my personal assistant on this number.
Goal: take my voice/text notes and reminders, and at ■time■ send me tomorrow's
list.
Tone: how I speak — short, no emoji.
NEVER: message anyone except me, never delete anything, never share my notes.
If a request is unclear → ask me one question.
```

WRITE YOUR OWN PROMPT

Session 2 homework · Fill in the blanks — this becomes your agent's job description

```
You are ■name your agent■, the assistant for ■business/person■.
Goal: ■trigger■ + ■action■ + ■done-condition■.
Tone: ■3 adjectives■. Language: ■which■.
Tools allowed: ■1■, ■2■, ■3■ (max).
NEVER: 1) ■hardest rule■ 2) ■second rule■ 3) ■third rule■.
If asked for something on the NEVER list → use ask_human().
Log every action with log_sheet().
```

Loop pseudocode — copy the pattern

```
while not goal_done and steps < 10:
    step = brain.think(goal, history)
    result = act(step.tool, step.inputs)
    history.append(observe(result))
    if result looks wrong:
        history.append(ask_human(result))
```

Adapt it to YOUR agent below:

HOMEWORK — SESSION 1

Your 3 wish-tasks · Due: before Session 2

Rules: make each task concrete — WHO it is for, WHAT it does, WHEN it happens, HOW OFTEN. If you cannot say when it is done, it is not a task yet.

Wish-task #1

Who is it for?	
What should the agent do?	
When / how often?	
What does 'done' look like?	
Tools it would need	
What must it NEVER do?	

Wish-task #2

Who is it for?	
What should the agent do?	
When / how often?	
What does 'done' look like?	
Tools it would need	
What must it NEVER do?	

Wish-task #3

Who is it for?	
What should the agent do?	
When / how often?	
What does 'done' look like?	
Tools it would need	
What must it NEVER do?	

HOMEWORK — SESSION 2

Design wish-task #1 · Due: before Session 3

GOAL (trigger + action + done)	
TRIGGER — what wakes the agent?	
TOOLS (name, inputs, output)	
STOP CONDITION (max steps / time)	
ESCALATE when… (ask_human)	

My agent's instruction prompt (8–12 lines)

Check yourself: does the prompt say WHO, WHAT GOAL, TONE, and at least one hard NEVER? If not, add it.

HOMEWORK — SESSION 3

Put it on WhatsApp · Due: before Session 4

Agent's NAME	
First message it sends (identity intro, 2 lines)	
Guardrails: NEVERs, caps, human gates	
Rate limits (msgs/min, night mode?)	
What needs MY approval before it acts?	

Sketch the round trip for YOUR task

Customer message → WhatsApp Web → MCP bridge → agent loop → reply → log. Draw arrows and label what happens at each step for your task:

HOMEWORK — SESSION 4

Ship it & earn · Due: Day 30

Shipping checklist — tick every box:

Restart-proof	Agent restarts itself after a crash
Logs	Writes logs I actually read weekly
Guardrails ON	Step caps + human gates still active
One-command pause	I can stop it instantly
Friend-tested	Someone else tried it and it held
Certificate submission	Goal + tools list + prompt + 1 screenshot

First client plan

Client 1 (who?)	
Their repetitive task	
Hours it eats monthly	
My setup fee / monthly care price	
Demo date I will offer	

MY FIRST AGENT — MASTER WORKSHEET

One page to run your agent. Photocopy this for agent #2 and #3.

1. NAME	
2. GOAL	
3. TRIGGER	
4. TOOLS (max 4)	
5. PROMPT (paste final version)	
6. GUARDRAILS	
7. STOP CONDITION	
8. LOG LOCATION	
9. FIRST TEST DATE	
10. FIRST USER	

The 30-day map

Week 1 (S1)	■ anatomy learned ■ 3 wish-tasks written
Week 2 (S2)	■ goal written ■ tools listed ■ prompt drafted
Week 3 (S3)	■ trigger set ■ guardrails set ■ round trip sketched
Week 4 (S4)	■ shipped ■ first user ■ first invoice / first ask

SAFETY CARD — pin this above your desk

From Session 3 · The three walls

IDENTITY	It says it is an assistant — never a human. One name, everywhere, no borrowed photos, no fake feelings.
PACING	One reply per customer message. Night mode. Hard rate limit. No broadcasts to strangers, ever.
HUMAN APPROVAL	Payments, discounts, dates, promises, refunds → ask_human() first. No answer in 15 min → wait, do not guess.

Failure map — where it will break first

Loops forever	Add step cap + stop condition
Robotic replies	Rewrite prompt; add one worked example
Misses messages	Check WhatsApp Web login / bridge
Invents facts	Give it a tool that reads real data
Tool-call storms	One job per tool; stricter NEVER list

Notes from class

GOAL → THINK → ACT → OBSERVE → REPEAT

You started with a wish-list.
You leave with an agent.

Stay hungry. Stay foolish.

AGENT SCHOOL — COHORT 1